

Day 3: Simple and complex data types

1 Introduction

Programs have lots of data, they are basically ways of getting data, transforming data, and giving data back to the world. Data types, as we have already seen, tell you which kinds of data you have in your program.

Computers store data in their memory in the form of bits. The word “bit” is a contraction of “BInary digiT”, so a bit is either 1 or 0, true or false, high or low, on or off. Bits are organised in groups of 8 called *bytes*. These bytes —groups of eight bits— are used to store any data in the memory of the computer. When you have got 1024 bytes, you have got a kilobyte or kB; when you have 1024 kB you have got a megabyte or MB, and so on for gigabytes (GB), terabytes (TB), and petabytes (PB). Note that in computing everything is measured in powers of 2 ($2^3 = 8$, $2^{10} = 1024$) and not in powers of 10 as in normal life (10, 100, 1000...). That is because computers count with bits (2) and human beings¹ count with their fingers (10).

2 Simple data types

Simple data types can be thought as boxes in the computer’s memory. Every time you declare a variable in your program, you can think² of the computer as creating a little box in its memory to store your variable. That box has two tags on it: one of them holds the name of the variable and the other holds the type (Figure 1). In the same way, you can think of assignment as putting a value inside that box (Figure 2).


2.1 Static typing and dynamic typing

You have already seen that Groovy, as most languages (including Java and Java Decaf), puts some restrictions to what you can use as a “name” tag for your boxes. You know that they have to start with a letter: “count” is a valid name but “/me” or “5thNumber” are not. You also know that some words cannot be used by the programmer for variables names because it would be confusing for the computer, e.g. “while”, “if”, “System”, or “println”.

Many programming languages also place one restriction on the “type” tag: once you decide the type of a variable, you cannot change it. This is like people having static opinions that they do not want to change. This type of languages are called *statically typed languages* and, contrary to

¹Not all humans used only the ten fingers in their hands to count. Some languages like French or Irish have remains of base-20 counting.

²This is only a metaphor and is not supposed to be an accurate description of how memory is managed on a modern computer. Explaining how things like the registers, the stack, and the heap work, the differences between actual machines and virtual machines, etc; are out of the scope of this document.



```
int i;
int j = 1;
```





Figure 1: Declaring a variable can be seen as creating a box. The box has a tag for the name and another for the type.



```
i = 2;
j = 2;
```




Figure 2: Assigning a value to a variable can be seen as putting a value in the box. If there was something in the box, it is overwritten and lost forever.

people with fixed ideas, they are not necessarily bad or obnoxious. Java is an example of statically typed language.

Some programming languages allow you to change the type of your variables as you go along, so you can have a variable that sometimes is an int and later in time is a boolean or a String. These languages are called *dynamically typed languages*. They have pros and cons compared to their statically typed counterparts. Groovy is an example of a dynamically typed language. The following code excerpt is valid in Groovy but not in Java. In Java a variable cannot be a String, then become a boolean, and then a String again. Note that *true* is a boolean while “true” (in quotes) is a String.

```
String str = "This is a string";
str = true;    // This is different to: str = "true"
if (str) {
    str = "This is a different string";
}
```

There are many more things to know about typing in programming languages (strong vs. weak, inferred vs. manifest, duck typing, and much more) but for now it will not be necessary to go into those details.

2.2 Most common simple types

2.2.1 Integer numbers

Integers (**int**) are probably the most used simple data type, as integer numbers are used for two of the most common operations in computing: counting and indexing. Integers use 32 bits of memory, and an integer variable can hold values between -2,147,483,648 and 2,147,483,647 (inclusive). This data type is large enough for the numbers needed in 90% of programs most people write. We have already seen how to use it:

```
int count = 1
```

There are two other kinds of integers for some special uses. When a program needs very large (positive or negative) values, there is a type **long**, for *long integer*. It is called “long” because it uses 64 bits instead of 32, which means a long integer variable can hold values between $-9.22 \cdot 10^{18}$ and $9.22 \cdot 10^{18}$. There is also a “short” integer that uses only 16 bits of memory (**short**); this was sometimes useful to save memory when Java was first released in 1995 and computers were more limited but is hardly ever used with modern computers.

2.2.2 Floating-point (decimal/rational) numbers

Not all numbers are integer. Examples of common non-integer numbers include the result of a division of two integers where one is not a multiple of the other, and real-measurements like your height, your weight, and the distance between your workplace and your home (unless you work at home). In maths, these are called real numbers. In computing they are usually called floating-point numbers and are represented as a list of significant numbers and an exponent. The term “floating-point” refers to the fact that the decimal point can “float”, that is, it can be placed anywhere in the number as long as the the exponent is changed accordingly.

$$1.23 \cdot 10^{-3} = 123 \cdot 10^{-5} = 0.00123 \cdot 10^0$$

We cannot use superindex notation in a plain-text file, so we need a special way of writing these numbers. Those three ways of depicting the number above can be written as `1.23E-3`, `123E-5`, and `0.00123`, and the three are equivalent.

In Groovy and Java real numbers are usually represented with the simple data type `double`, that uses 64 bits. There is also a 32-bit version called `float` but, as with `short`, it is hardly ever used today.

Important note. Floating-point numbers do *not* have infinite precision, and operating with them can cause rounding errors. There is a special type for representing real values where precision is paramount (like in banking): `BigDecimal`.

Equality. Due to rounding errors, it does not make sense to test for equality among real numbers as we can do with integers. Two numbers could be notionally the same but be different due to rounding errors, so they are never compared with equality with `==`. Instead, what is usually done is test whether the difference is less than some precision limit appropriate for the application (e.g. `1.0E-6`), as in the example below. Note: `Math.abs()` returns the absolute value of the number inside the brackets, e.g. `Math.abs(-3)` returns 3.

```
double d1 = 2 / 3
double d2 = 2 * 1000 / 9000 * 3
// WRONG!
if (d1 == d2) {
    // this is not printed due to rounding errors
    println "They are the same (wrong comparison)"
}
// RIGHT!
if (Math.abs(d1 - d2) < 10E-6) {
    println "They are the same (right comparison)"
}
```

2.2.3 Boolean (binary) values

This simple data type represents one bit of information. It can hold the values `true` and `false`.

2.2.4 Characters

Text is composed of characters: `'a'`, `'b'`, `'c'`... The `char` simple type is used to represent characters. It uses 16 bits, meaning it can represent any of 65,536 different characters.

Actually, the 16 bits of a `char` represent a Unicode symbol. Unicode is a computing industry standard for the consistent encoding, representation and handling of text expressed in most of the world's writing systems. It includes symbols from most writing systems in the world, including alphabets like Latin, Cyrillic, Arabic, or Hebrew; syllabaries like Japanese katakana and hiragana, or Cherokee; and many more.

You may have noticed that we have not mentioned `String` yet. This is because `String` is a complex type.

3 Complex types

Complex types are types of data that do not fit in a box, not even in one of the big 64-bit boxes used for `double`. Because they do not fit in the “boxes”, computers have to store them somewhere else. However, they also need to know where they are... and that is what the boxes are used for.

Modern computers have *a lot* of memory. Long forgotten are the days when Bill Gates said: “640kB of memory should be enough for everything”. Part of a computer’s memory is used for the boxes (in a part of memory called “the stack”) and most of the rest is used for everything else, including complex data (that part is called, quite unceremoniously, “the heap”).

When your Groovy, Java Decaf, or Java code uses some complex data, the computer stores that data in some region of the heap —identified by a *memory address*; then it stores the address in a box in the stack, much in the way it stores integers and booleans. This looks similar to Figure 3. The memory address in the box can be seen as “pointing to” the place in memory where the real data is stored. For this reason, we will call it a *pointer*.




Figure 3: A String is a type of complex data type. The data itself is stored in the heap and its address is stored in box in the stack —pointing to it.

Using complex data types is different from using simple types in several ways. For starters, as complex types are not stored in the stack but in the heap, you have to *allocate* memory-space in the heap to store the data. This is done with the reserved word `new`; you also need brackets to pass arguments (as if using a method: this is because complex types have something called a *constructor method* that we will study in more detail later). We are going to see now several examples, starting with the pervasive Strings.

One final note about names. Complex data types usually have names that start with capital letters. This is not compulsory but it is what everybody expects: simple types with non-capital letters (e.g. `int`, `double`) and complex letters with capital letters (e.g. `String`, `Array`, `List`, `Customer`...). If you do not capitalise your complex types other people will get very confused when they read your code. This is not only impolite, it is also unprofessional.

3.1 Declaration and initialisation

When working with simple types, the memory is always used as soon as you declare the variable. In other words, your program will use the same amount of memory if you type `int i` and if you type `int i = 1`: it always uses 32 bits of your computer’s memory.

This is not true with complex types. When a complex type is declared (e.g. `String str`) the computer only reserves the box for the pointer, but nothing else. It is only when the variable is *allocated* (e.g. `str = new String()` or `str = new String("This is a String")`) that the total

amount of memory is used. Note that the difference is huge. Boxes are 32-bit or 64-bits long, but there is no limit to the size of a complex type: it could be several kB or even MB (types so big are a rarity, though).

Sometimes the pointer will not point to any address in memory. This can be useful in some cases that will become clearer as we learn more about programming, including error detection and release of memory that is no longer needed (remember that complex types can use a lot of memory). To make a pointer point to nowhere, we use the reserved word `null` (with non-capital letters).

```
String str
str = null
```

A pointer pointing to null is called a *null pointer*. This basically means that the address in the box is zero (rather than an obscure hexadecimal number like `0x1a3ec74`); the computer knows there is never anything at address zero, so it knows a null pointer is not being used to point to any real data (Figure 4).




Figure 4: A null pointer is basically a zero address, pointing nowhere.

As a null pointer does not point to any real data, if you try to access a variable that is pointing to null, the computer will complain with a `NullPointerException`, as in the following example (try it!):

```
String str = null
println str.length()
```

3.2 String

Strings are everywhere. Every program, except the most trivial, uses strings: user names, passwords, addresses, configuration options, data input from the keyboard, a webpage read through the Wi-Fi connection... almost anything is a `String`. It is the most widely used complex type.

At their most basic, strings are sequences of characters. You already know how to read a piece of text from the user:

In Groovy: `String str = System.console().readLine()`

In Java Decaf: `String str = nextLine()`

If you want to create a string in your program without reading it from the user, the basic form of creating a string is like this:

```
String str = new String("This is a String")
```

but you already know that this can be made in an easier way:

```
String str = "This is a String"
```

The former two statements are equivalent in Groovy³, only the second is more convenient and most people use it instead of the other.

It is important to note that a String with only one character is still a String.

You already know that you can do some things with strings using *methods*. Although we are going to learn more about methods on the next section, we can introduce some of them now to help with the exercises. Assuming you have a String called `str`, you can use the following methods:

str.length(): this method returns the length of the string.

str.charAt(): this method requires an integer inside the brackets, and returns the character at the specified position; note that the first character is at position 0, not 1.

str.substring(): this method requires two integers inside the brackets, separated by a comma, and returns another string that begins at the character specified by the first number and extends to the character specified by the second number minus one⁴.

Here is a brief example of how they are used. Note that if you want to use double quotes inside double quotes, you have to *escape* them using the backslash (“\”) sign: the backslash will make Groovy and Java treat the next double quote as a literal double quote and not as the end of the string.

```
String str;
str = new String("This is an example")
println "Initial string: \"" + str + "\"" // Note the escaped quotes!
int l = str.length()
println "The length of the string is " + l
char c = str.charAt(0)
println "The first character is " + c
String str2 = str.substring(8,18)
println "The substring from char 8 to char 17 is \"" + str2 + "\""
```

The result of this little program is:

```
Initial string: "This is an example"
The length of the string is 18
The first character is T
The substring from char 8 to char 17 is "an example"
```

³They are NOT equivalent in Java, but the difference is subtle and we will learn more about it when the time comes.

⁴This results in the substring having a length of `secondNumber - firstNumber`.

3.3 You own structures: classes

String is the most common complex type, but it is not the only one. As a matter of fact, programmers can create their own complex types very easily. In order to create new types of complex data, we use the keyword `class`. You can think of it as creating new classes of data.

A new class of data must have a name (remember: by convention complex data types start with capital letters, classes too) and be defined in between curly brackets. A complex type is composed of several other types, simple or complex. Let's see an example:

```
class Person {
    String name;
    int age;
}
```

This code defines a new class of data that represents a person, and it is composed of a String for the name and an integer for the age. As you can see, both simple and complex types can be used inside a class. You can even have data on a specific class inside the same class, as in the extended example below:

```
class Person {
    String name;
    int age;
    Person father;
    Person mother;
}
```

Now a person has a String for a name, an integer for the age, and two variables of type Person for the father and the mother. Once you start working with complex data on a regular basis (which means from now on) you can see that it is crucial to use good identifiers for your names. Compare the former class with this very-badly-named example:

```
class P {
    String s;
    int n;
    P p1;
    P p2;
}
```

Both examples contain the same information at a logical level, and the computer will behave equally well with both. However, human programmers reading the code will have a hell of a time trying to understand what the second class really represents, while the first example is obvious. Remember: **source is written once but is read many times**, so your code should be as clear as possible.

Variables in a class are often called *fields*. Sometimes they are called *member fields* because they can be seen as being “part of” (members of) the class. So the class `Person` can be said to have four fields: one of type `String`, one of type `int`, and two of type `Person`. Fields in a class are accessed using a dot, as in the following example:


```

Person employee = new Person();
employee.name = "John Smith";
employee.age = 45;
println "BOSS: How old are you, " + employee.name + "?"
println "EMPLOYEE: I am " + employee.age + " years old.";

```

In this example we have created one variable of class `Person`, a so-called *instance* of the class. Each of these instances of a class are also called *objects*. Every time you use the keyword **new** you are creating another object (this involves looking for a place in memory to store it and place the corresponding pointer pointing to it). You can think of a *class* as a blueprint, and thus creating a new instance (by using **new**) is creating a new object according to that blueprint.

Note that fields on an object are accessed using a dot to separate the name of the variable that identifies the object (e.g. `employee`) and the name of the variable inside the object, the field (e.g. `name`), as in `employee.name`.

We can make the example slightly more complicated by using more than one instance of the class `Person`.

```

Person john = new Person();
john.name = "John Smith";
john.age = 35;
Person mary = new Person();
mary.name = "Mary Smith";
mary.age = 32;
Person student = new Person();
student.name = "John Smith, Jr.";
student.age = 5;
student.father = john
student.mother = mary
println "TEACHER: How old are you, " + student.name + "?"
println "LITTLE JOHN: I am " + student.age + " years old, sir.";
println "TEACHER: Who is your mother?"
println "LITTLE JOHN: " + student.mother.name + ", sir.";

```

There are two important things to learn from this example. First, we can access data inside an object that is inside another object. Actually, there is no limit to the level of encapsulation you can achieve with complex data (as long as you do not run out of memory). If you want to access a field of a field you only need to use two dots, as in `student.mother.name`. Second, we have not initialised all fields in the objects of class `Person`. When you create a new object, the fields of the object have to be initialised as any other variable, otherwise they will not have the values you want them to have. In this example, we do not know who are the father and the mother of objects `john` and `mary`. If we try to access them, the computer will complain. On the other hand, if the mother of object `mary` was properly initialised, we could know the name of the grandfather of the student by typing something like:

```

println "The father of my mother is " + student.mother.father.name;

```

Next day we will see how to add methods to your classes, like those that you know from `String`.

3.4 Boxed types: Integer, Double, and Character

For every simple type in Groovy, Java Decaf, and Java, there is a “boxed” complex version. The most important ones are summarised in Table 1. Note that the name of the complex type uses capital letters and is sometimes a longer, complete-word version of the simple type’s name.

| Simple | Complex |
|--------|-----------|
| int | Integer |
| double | Double |
| char | Character |

Table 1: Most important simple types and their boxed counterparts

The complex version of a simple type works mostly like any other complex type. The variable itself contains a pointer that points to the address in memory where the actual value is stored (Figure 5). Why have complex versions of simple types? Other data types have a lot of information (i.e. many characters in a String) and that is why they need to be stored indirectly, but this is not the case with boxed types. Why then? Remember that complex types of data can have their own methods, and sometimes this is very useful. You have already come across one of the most used ones: `Integer.parseInt()`, which is used to convert strings that contain only a number into an `int`. Boxed types have several other methods, and we will see some of them, but `parseInt()` is the one more frequently used.




Figure 5: A boxed type, like any complex type, consists of a pointer to a place in memory where the actual data is stored

Why are these complex types called “boxed”? This is because they can be seen as a box wrapping around a simple type. In the old days, using boxed types was cumbersome and boring. You could need to type something like:

```
int i = 1;
Integer boxedI = new Integer(i);
int i2 = boxedI.intValue()
```

You had to explicitly create the boxed type like any other complex type by using `new`, and then you needed a method of the boxed type to get the simple type from inside. This was a lot of work when you had many variables and had to perform conversion from simple type to boxed type very often, for example if you have to make comparisons of data between two different sources (one simple type, the other boxed type). This why Java (and Groovy) now allow what is called *auto-boxing*, which is just a fancy name to say that you can use them in exactly the same way and

the computer takes care of all the boxing-in and boxing-out. In other words, the following code is perfectly legal:

```
int i1 = 1
Integer i2 = 1
if (i1 == i2) {
    println "Nowadays you can use boxed and simple types together!"
}
```

Note that you do not need to write explicitly `i2 = new Integer(1)`, but `i2` is still a complex boxed type, its box contains just a pointer to some place in memory where a “1” is stored. The variable `i1`, on the other hand, is a *bona fide* integer variable, a box with a “1” inside (see Figure 6). However, thanks to auto-boxing, you can compare one with the other: the computer will automatically extract the “1” inside the boxed type to perform the comparison with the other “1”. You do not need to do it yourself!




Figure 6: Even if `i1` and `i2` are, strictly speaking, two very different things, auto-boxing makes it easy to work with both types at the same time.

3.5 A final note on terminology

If you read other books or web pages about Groovy or Java, you may notice that they use terms that are different from the ones we have used in this section. We go through some of them here for the sake of clarity.

Simple type: A type that is stored in its own box, usually 32-bit or 64-bit long. These are usually called *primitive* types in the Java world. I prefer to call them simple data types to distinguish them clearly from complex data types.

Complex type: A type that has two parts: the box contains a pointer, and it points to a place in memory where the actual value is stored. In the Java world, these are simply called *classes* and *objects*, which is precise but fails to be explicit on the difference between simple and complex data types, and this can be a source of confusion.

Pointer: The content of the box in a complex type, a memory address where actual data is stored. This is sometimes called a *reference* or a *handle* to prevent confusion with pointers of

other languages (like C) that behave in a slightly different way. I think many programming languages have constructs that share the same name and behave in slightly different ways (Strings is a major example) so this is not of much concern. Besides, it looks to me like very confusing to access a null “reference” or a null “handle” and get a `NullPointerException` as a consequence.